

OSPP

Threads

# Advantages of Threads

- The overhead for creating a thread is significantly less than that for creating a process
- Multitasking, i.e., one process serves multiple clients
- Switching between threads requires the OS to do much less work than switching between processes

# Drawbacks of Threads

- Not as widely available as longer established features
- Writing multithreaded programs require more careful thought
- More difficult to debug than single threaded programs
- For single processor machines, creating several threads in a program may not necessarily produce an increase in performance

# main thread

- initial thread created when `main()` is invoked by the process loader
- once in the `main()`, the application has the ability to create daughter threads
- if the main thread returns, the process terminates even if there are running threads in that process, unless special precautions are taken
- to explicitly avoid terminating the entire process, use `pthread_exit()`

# Create thread

- `int pthread_create( pthread_t *thread, pthread_attr_t *attr, void *(*thread_function)(void *), void *arg );`
- 1st arg – pointer to the identifier of the created thread
- 2nd arg – thread attributes. If null, then the thread is created with default attributes
- 3rd arg – pointer to the function the thread will execute
- 4th arg – the argument of the executed function
- returns 0 for success

# Waiting threads

- `int pthread_join( pthread_t thread, void **thread_return )`
- main thread will wait for daughter thread *thread* to finish
- 1st arg – the thread to wait for
- 2nd arg – pointer to a pointer to the return value from the thread
- returns 0 for success
- threads should always be joined; otherwise, a thread might keep on running even when the main thread has already terminated

# Sample Pthreads Program in C

- The program in C calls the pthread.h header file. Pthreads related statements are preceded by the pthread\_ prefix (except for semaphores).
- How to compile:
  - gcc hello.c -pthread -o hello

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <unistd.h>
```

```
int main(int argc, char **argv){
    pthread_t t1;
    int thread_id = 1;

    if ( (pthread_create(&t1, NULL, (void *)&worker, (void *)&thread_id)) != 0) {
        printf("Error creating thread\n");
        exit(1);
    }

    pthread_join(t1, NULL);

    return 0;
}
```

```
void worker(void *a) {
    int *cnt = (int *)a;

    printf("This is thread %d\n", *cnt);
    pthread_exit(0);
}
```



# Thread Synchronization Mechanisms

- Mutual exclusion (mutex):
  - guard against multiple threads modifying the same shared data simultaneously
  - provides locking/unlocking critical code sections where shared data is modified
  - each thread waits for the mutex to be unlocked (by the thread who locked it) before performing the code section

# Basic Mutex Functions

```
int pthread_mutex_init(pthread_mutex_t *mutex, const
    pthread_mutexattr_t *mutexattr);
int pthread_mutex_lock(pthread_mutex_t *mutex);
int pthread_mutex_unlock(pthread_mutex_t *mutex);
int pthread_mutex_destroy(pthread_mutex_t *mutex);
```

- a new data type named `pthread_mutex_t` is designated for mutexes
- a mutex is like a key (to access the code section) that is handed to only one thread at a time
- the attribute of a mutex can be controlled by using the `pthread_mutex_init()` function
- the lock/unlock functions work in tandem

```

#include <pthread.h>
...
pthread_mutex_t my_mutex;
...
int main()
{
    int tmp;
    ...
    // initialize the mutex
    tmp = pthread_mutex_init( &my_mutex, NULL );
    ...
    // create threads
    ...
    pthread_mutex_lock( &my_mutex );
        do_something_private();
    pthread_mutex_unlock( &my_mutex );
    ...
    ...
    pthread_mutex_destroy(&my_mutex );
    return 0;
}

```

- Whenever a thread reaches the lock/unlock block, it first determines if the mutex is locked. If so, it waits until it is unlocked. Otherwise, it takes the mutex, locks the succeeding code, then frees the mutex and unlocks the code when it's done.

# Semaphores

- Counting Semaphores:
  - permit a limited number of threads to execute a section of the code
  - similar to mutexes
  - should include the `semaphore.h` header file
  - semaphore functions have `sem_` prefixes

# Basic Semaphore Functions

- creating a semaphore:

```
int sem_init(sem_t *sem, int pshared, unsigned int value);
```

- initializes a semaphore object pointed to by sem
- pshared is a sharing option; a value of 0 means the semaphore is local to the calling process
- gives an initial value value to the semaphore

- terminating a semaphore:

```
int sem_destroy(sem_t *sem);
```

- frees the resources allocated to the semaphore sem
- usually called after pthread\_join()
- an error will occur if a semaphore is destroyed for which a thread is waiting

# Basic Semaphore Functions

- semaphore control:

```
int sem_post(sem_t *sem);
```

```
int sem_wait(sem_t *sem);
```

- `sem_post` *atomically* increases the value of a semaphore by 1, i.e., when 2 threads call `sem_post` simultaneously, the semaphore's value will also be increased by 2 (there are 2 atoms calling)
- `sem_wait` *atomically* decreases the value of a semaphore by 1; but always waits until the semaphore has a non-zero value first

```

#include <pthread.h>
#include <semaphore.h>
...
void *thread_function( void *arg );
...
sem_t semaphore;    // also a global variable just like mutexes
...
int main()
{
    int tmp;
    ...
    // initialize the semaphore
    tmp = sem_init( &semaphore, 0, 0 );
    ...
    // create threads
    pthread_create( &thread[i], NULL, thread_function, NULL );
    ...
    while ( still_has_something_to_do() )
    {
        sem_post( &semaphore );
        ...
    }
    ...
    pthread_join( thread[i], NULL );
    sem_destroy( &semaphore );
    return 0;
}

```

```
void *thread_function( void *arg )  
{  
    sem_wait( &semaphore );  
    perform_task_when_sem_open();  
    ...  
    pthread_exit( NULL );  
}
```

- the main thread increments the semaphore's count value in the while loop
- the threads wait until the semaphore's count value is non-zero before performing `perform_task_when_sem_open()` and further
- daughter thread activities stop only when `pthread_join()` is called